

Kaggle - Zillow Home Prices EDA

Lilian Cheung

March 18, 2018

```
runthis <- FALSE

if(runthis==TRUE) {
  path_properties_2016 <- "C:/Lilian/Documents/Technical Training/R Projects/Kaggle/Zillow Home Prices/properties_2016/properties_2016.csv"

  path_properties_2017 <- "C:/Lilian/Documents/Technical Training/R Projects/Kaggle/Zillow Home Prices/properties_2017/properties_2017.csv"

  path_train_2016 <- "C:/Lilian/Documents/Technical Training/R Projects/Kaggle/Zillow Home Prices/train_2016_v2/train_2016_v2.csv"

  path_train_2017 <- "C:/Lilian/Documents/Technical Training/R Projects/Kaggle/Zillow Home Prices/train_2017/train_2017.csv"

  ## Training Data - 2016 (properties and logerrors of sales)
  properties_2016 <- fread( file=path_properties_2016, header=TRUE )
  train_2016 <- fread(file=path_train_2016, header=TRUE)

  ## Training Data - 2017 (properties and logerrors of sales)
  properties_2017 <- fread( file=path_properties_2017, header=TRUE )
  train_2017 <- fread(file=path_train_2017, header=TRUE)
}
```

```
runthis <- FALSE

if (runthis==TRUE){
  # Merge training data from 2016
  training_2016 <- merge( properties_2016, train_2016, by="parcelid" )

  # Merge training data from 2017
  training_2017 <- merge( properties_2017, train_2017, by="parcelid")

  # Combine 2016 and 2017 data into a dataframe
  training_all <- rbind(training_2016, training_2017)

  # Save Locally
  save_path <- "C:/Lilian/Documents/Technical Training/R Projects/Kaggle/Zillow Home Prices/training_all.csv"
  write.csv( training_all, file=save_path, row.names=FALSE)
}

# Read in dataset
training_all_path <- "C:/Lilian/Documents/Technical Training/R Projects/Kaggle/Zillow Home Prices/training_all.csv"
training_all <- read.csv( file=training_all_path, header=TRUE)

## Sample some percent of dataset for EDA
set.seed(123)
percent <- 0.10
training_sample <- sample_n( training_all, percent*nrow(training_all), replace=FALSE )

## remaining data
# remaining_data <- dplyr::filter( training_all, !parcelid %in% training_sample[["parcelid"]] )

## Save Locally
save_path <- "C:/Lilian/Documents/Technical Training/R Projects/Kaggle/Zillow Home Prices/training_sample.csv"
write.csv( training_sample, file=save_path, row.names=FALSE)
```

```
## Read in training sample
read_path <- "C:/Lilian/Documents/Technical Training/R Projects/Kaggle/Zillow Home Prices/training_sample.csv"
training_sample <- read.csv( file=read_path, header=TRUE)
training_sample <- mutate(training_sample, transactiondate = as.Date(transactiondate))

## Reorder training sample by date
training_sample <- arrange( training_sample, transactiondate )
```


```
## Explore transaction counts in sample

# Number of transactions per day
counts_transactions <- count( training_sample, transactiondate )

n_transactions_plot <- ggplot( counts_transactions, aes(x = transactiondate, y = n) ) +
  geom_line( color = "skyblue" ) +
  xlab("Transaction Date") +
  ylab("Number of Transactions") +
  ggtitle("Number of Transactions Per Day")

print( n_transactions_plot )
```

Number of Transactions Per Day



Number of transactions per month in our dataset (note that October 2016 dips since part of the dataset is in Zillow's test dataset)

```
dt_training_sample <- data.table( training_sample )
```

```
counts_transactions_month <- dt_training_sample[ , .N, by = .(year(transactiondate), month(transactiondate)) ]
```


Graph (uses xts package)

Manipulate data into time series

```
transactions_month <- counts_transactions_month[["N"]]
```

```
transactions_month <- ts( counts_transactions_month[["N"]], frequency = 12, start = 2016)
```

```
plot(as.xts(transactions_month), major.format = "%Y-%m", title="Number of Transactions per Month")
```



```
## Exploring the continuous predictors
```

```
## Convert variable types to factor as necessary
```

```
to_factor <- c( "airconditioningtypeid", "architecturalstyletypeid",  
  "buildingqualitytypeid", "buildingclasstypeid", "decktypeid",  
  "fips", "heatingorsystemtypeid",  
  "parcelid", "poolcnt", "pooltypeid10", "pooltypeid2", "pooltypeid7",  
  "propertylandusetypeid", "rawcensustractandblock",  
  "censustractandblock",  
  "regionidcounty", "regionidcity", "regionidzip",  
  "regionidneighborhood", "storytypeid", "typeconstructiontypeid",  
  "assessmentyear", "taxdelinquencyyear")  
training_sample[ to_factor ] <- lapply( training_sample[to_factor], factor)
```

```
## Examine continuous variables
```


```
num_columns <- sapply(training_sample, is.numeric)  
cts_training_sample <- training_sample[ num_columns ]
```

```
# Correlogram
```

```
cor_matrix <- cor( cts_training_sample, use='pairwise.complete.obs' )
```

```
## Warning in cor(cts_training_sample, use = "pairwise.complete.obs"): the  
## standard deviation is zero
```


```
corrplot(cor_matrix, type='lower')
```



Explore Logerror, the response variable

```
logerror_histogram <- ggplot( training_sample, aes(x = logerror) ) +
  geom_histogram( bins = 20)

logerror_histogram
```




```
## Plot of Logerror per day
mean_logerror_daily <- dt_training_sample[, mean(logerror), by = transactiondate]
names(mean_logerror_daily) <- c("transactiondate", "mean_daily_logerror")

daily_mean_logerror_plot <- ggplot( mean_logerror_daily, aes(x = transactiondate, y = mean_daily_logerror) ) +
  geom_line( color = "darkred" ) +
  xlab("Transaction Date") +
  ylab("Mean Logerror") +
  ggtitle("Mean Logerror Per Day")


daily_mean_logerror_plot
```


Mean Logerror Per Day



```
## Does average logerror per day exhibit autocorrelation? No.  
acf( mean_logerror_daily[["mean_daily_logerror"]] )
```

Series mean_logerror_daily[["mean_daily_logerror"]]



```
## As expected, number of daily transactions does not have a high correlation with mean daily logerrors  
cor(x=mean_logerror_daily$mean_daily_logerror, y=counts_transactions$n)
```

```
## [1] 0.05175558
```

```
# 0.05
```

```
## Correlations between high logerrors and continuous features


## Arbitrarily define a high logerror as one that meets or exceeds 0.2
high_logerror_sample <- dplyr::filter( training_sample, abs(logerror) >= 0.25 )

## Correlations between continuous variables in the high logerror sample
num_columns <- sapply( high_logerror_sample, is.numeric)
cts_high_logerror_sample <- high_logerror_sample[ num_columns ]


# Correlogram
cor_matrix <- cor( cts_high_logerror_sample, use='pairwise.complete.obs' )
```

```
## Warning in cor(cts_high_logerror_sample, use = "pairwise.complete.obs"):  
## the standard deviation is zero
```

```
corrplot(cor_matrix, type='lower')
```



```
## A closer Look at columns of interest -- turns out most of these are missing, so this is not particularly helpful
cols_of_interest <- c( "poolsizesum", "yardbuildingsqft17", "finishedsquarefeet13", "finishedsquarefeet6", "logerror")
cor_matrix <- cor( cts_high_logerror_sample[, cols_of_interest], use='pairwise.complete.obs')
corrplot(cor_matrix, type='lower')
```



```
# poolsize: total square footage of all pools on property
# yardbuildingsqft: Patio in yard
# finishedsquarefeet13: Perimeter Living area
# finishedsquarefeet6: Base unfinished and finished area
```

```
## These variables are mostly missing, so are not very helpful - e.g.
# table( is.na(cts_high_logerror_sample$poolsize) )
```

```
## ANOVA matrix?
```